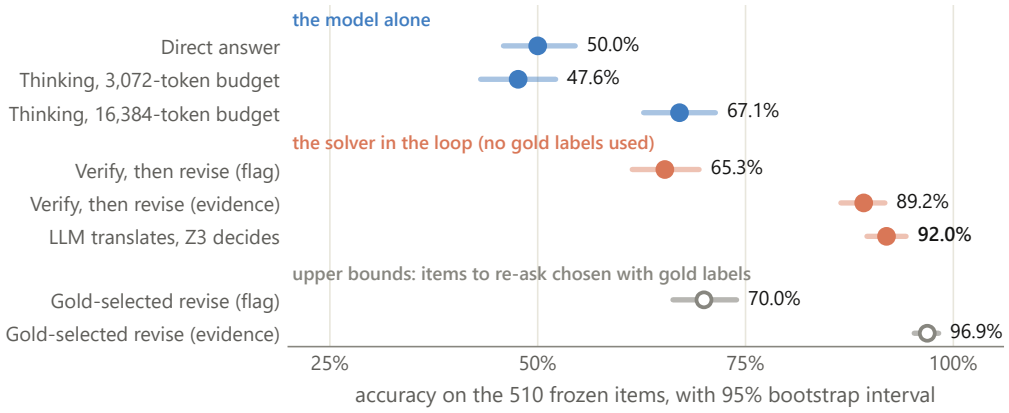
the model alone
Direct answer
50.0%
Thinking, 3,072-token budget
47.6%
Thinking, 16,384-token budget
67.1%
the solver in the loop (no gold labels used)
Verify, then revise (flag)
65.3%
Verify, then revise (evidence)
89.2%
LLM translates, Z3 decides
92.0%
upper bounds: items to re-ask chosen with gold labels
Gold-selected revise (flag)
70.0%
Gold-selected revise (evidence)
96.9%
25%
50%
75%
100%
accuracy on the 510 frozen items, with 95% bootstrap interval